



Intent becomes action.

Vision & Technical Architecture

Prepared for: Internal Team Review · Confidential

Founding Team

Maaz	Co-Founder & CTO — Pakistan
Gilson	Co-Founder & GTM — Curaçao
Nekha	Founding Engineer — Bay Area, US
Ashish	Founding Designer — India
Samantha	Lead Designer — Australia

1. The Problem With the Internet Today

The internet is the most powerful resource ever built. But accessing it still requires the user to do all the work. Search. Filter. Tab-switch. Compare. Paste. Repeat. Every platform optimises for its own engagement — not for the user's outcome. iAM is the first platform designed entirely around the user's intent.

Think about what it takes to plan a birthday trip for a friend. You open Google Flights in one tab, Booking.com in another, TripAdvisor in a third, Google Maps in a fourth, OpenTable in a fifth. You compare. You paste links into a group chat. You wait for people to reply. You lose the tab with the price you saw. You start over. You pay on four different platforms with four different accounts. Then you do the same for their gift. Then for what to pack. Then for the weather.

This is the current state of the internet. Seventeen platforms. Zero coordination. No memory. No action.

iAM replaces all of that with one sentence.

The user types what they want in plain English. iAM reads the intent, assembles every relevant result from across the internet, personalises it to who that person is, enables collaboration with anyone they tag, and completes every action — booking, purchasing, scheduling, gifting — in one place. The internet finally comes to the user.

2. What iAM Is

2.1 The Stage — the internet, assembled around you

The Stage is iAM's core surface. Not a search results page. Not a dashboard. A living, streaming workspace that assembles around a single prompt — pulling flights, hotels, restaurants, experiences, products, services, content, weather, maps, and more in parallel, then ranking everything against the specific user's history and preferences.

Every row in the Stage is a category of intent. Every card in a row is an actionable result — bookable, purchasable, reservable, or readable without leaving iAM. When the user is done, one checkout handles everything regardless of how many vendors are involved.

The Stage in one sentence: Type what you want in plain English. iAM builds the internet's answer around you, personalised, collaborative, and fully actionable — in under few seconds.

2.2 How iAM covers the entire internet


iAM does not crawl or host the internet's content. Instead, it is connected to the internet through a comprehensive and expanding layer of third-party APIs — one for every domain of human intent. When a user types a prompt, iAM determines which services are relevant, calls them simultaneously, and assembles the results into a unified Stage. The user never knows — or in some cases they need to know — which API answered.

DOMAIN	PROVIDER / APPROACH	WHAT IT DOES IN THE STAGE
Flights	Duffel	Real-time routes, pricing, seat selection — fully bookable
Hotels & stays	Duffel	Accommodations filtered by group size, dates, budget signal
Rental cars	Duffel	Category filters, passenger count, instant booking
Activities & tours	Viator	Museum tickets, river cruises, guided tours — bookable
Restaurant bookings	OpenTable	Availability matched to group size, date, cuisine preference
Weather	OpenWeatherMap	Local forecast for trip dates — informs packing and planning
Maps & discovery	Google Maps	Nearby POIs, malls, landmarks as a visual discovery row
Products & shopping	Retailer APIs	Direct product search and purchase across major retailers
Digital services	Freelance APIs	Designers, developers, copywriters, domain registrars

Home services	Booking APIs	Plumbers, electricians, cleaners, handymen
Health services	Clinic APIs	Doctors, dentists, therapists — appointment booking via Genie
Appointments	Calendly / Cal	Any service with scheduling — Genie books on your behalf
Payments	Stripe Connect	Single checkout across all vendors — split handled automatically

This is not a fixed list. Every new category of human intent iAM enters is enabled by connecting the relevant API layer. The Stage architecture is domain-agnostic — iAM does not need to be rebuilt to enter a new vertical. It needs a new integration and the appropriate intent parsing logic.

2.3 Three actors. One system.



3. How the LLM Helps the User

The language model is not the product. It is the intelligence layer that makes the product possible. iAM uses the LLM to do three things no rule-based system can: understand what the user means, generate responses that feel genuinely personal, and decide which actions to take next.

3.1 Intent parsing — from words to structured action

Every prompt a user types goes through an LLM pipeline before anything else happens. The LLM does not do keyword extraction. It understands meaning. "Something warm and not too expensive for the four of us who like history" is not a query a search engine can handle. The LLM parses it into a structured Parsed Intent object that drives every downstream API call:

FIELD	HOW THE LLM EXTRACTS IT
destination	Inferred from context — or clarified by Genie if ambiguous
origin	Extracted from profile or prompt context
dates	Natural language date parsing — 'next month', 'early June', specific dates
participants	@handles resolved to user IDs — their intent graphs merged into the Stage
group size	Inferred from participant count or explicit mention
budget signal	Inferred from language — 'budget', 'mid-range', 'luxury', 'affordable'
activity type	Historical preference, culture, outdoor, food-led, nightlife
constraints	Dietary, accessibility, timing, style preferences

3.2 Personalised responses — the LLM + intent graph together

For every user-facing response — a profile query answer, a Genie suggestion, a Stage ranking justification — the LLM does not generate from scratch. It generates with the user's intent graph as context. The graph is serialised and injected into the LLM prompt alongside the user's question. This means the response is not what the LLM knows about the world in general — it is what the LLM knows about the world through the lens of this specific person.

Example: Two users both type "recommend a hotel in Tokyo." User A's intent graph shows premium travel, solo trips, and a preference for design hotels. User B's graph shows budget travel and group trips. The LLM receives the same question — and two entirely different answers, because the context it reasons from is different.

3.3 Language as the universal interface

The LLM makes it possible to interact with iAM the same way you interact with a knowledgeable person. You do not configure filters. You do not navigate menus. You do not learn the platform's logic. You say what you want, and the system figures out the rest. This is the shift iAM represents: from interface-centric computing to intent-centric computing.

4. The Intent Graph

The Intent Graph is iAM's memory, personalisation engine, and competitive moat — all in one. It is a continuously updated, per-user data structure that captures not just what someone bought, but who they are becoming. Every action on the platform makes the next Stage smarter.

4.1 What it stores

SIGNAL TYPE	WHAT IAM CAPTURES AND WHY
Destinations	Where the user has searched, browsed, locked, and booked — with recency weighting
Spending signal	Budget, mid-range, or premium — inferred silently from choices, never asked
Activity preferences	Museums, outdoor, nightlife, food-first — derived from engagement across Viator and OpenTable
Travel style	Solo, couple, group — inferred from participant counts across all sessions
Seasonal patterns	When the user tends to travel and to where — used by Genie for proactive prompts
Taste signals	Design aesthetics, brand preferences, cuisine preferences — inferred from browsing + locking
Outcome history	What was actually booked — the strongest signal for future ranking decisions
Service preferences	Preferred airlines, hotel chains, restaurant types — built over time

4.2 How it is built — technically

Structured fields

The base layer of the intent graph is a structured DB document. Every completed booking, every locked card, every prompt submitted writes signals back to this document. The writes are weighted — a completed booking carries far more signal than a browse. Signals decay over time with recency weighting so the graph reflects who the user is now, not who they were a year ago.

Vector embeddings

The graph is also embedded as a vector. Every significant interaction — a prompt, a document upload, a completed booking — is passed through a text embedding model. The resulting vectors are stored in a vector database (Pinecone or pgvector). This enables semantic similarity search across the user's full history — the foundation of how Genie retrieves context.

RAG retrieval — Retrieval-Augmented Generation

When Genie or a profile prompt needs to answer a question about a user, it does not load the entire intent graph. It runs a RAG pipeline: the incoming query is embedded, the vector store is searched for the most semantically relevant fragments of that user's history, and those fragments are injected into the LLM prompt as context. The LLM generates its response grounded in retrieved, real preference data — not hallucinated guesses.

How a profile query works technically: Visitor asks: "What's Nekha's favourite colour?" → Query embedded → Vector search over Nekha's intent graph → Relevant preference fragments retrieved → LLM prompt constructed: [system: answer about Nekha using only this context] + [retrieved fragments] + [user question] → Personalised answer returned.

4.3 How it drives the Stage

Before any API call fires in the Stage assembly, all participant intent graphs are fetched and merged into a shared StageContext. This merged context is passed to the relevance gate — every API result must score above a threshold on Intent Fit + User Fit before it is even considered for ranking. The ranking formula then orders the qualifying results:

FACTOR	WHAT IT MEANS
Intent Fit	How closely this result matches the semantic meaning of the user's prompt
User Fit	How well this result aligns with this user's intent graph history and preferences
Outcome History	How often this type of result has led to completed purchases for similar users
Bid	A commercial signal that can shift position within the qualified set

The moat that compounds: A user's Intent Graph is the most personalised AI model of their preferences that has ever existed. It cannot be replicated by any competitor who acquires that user without their history. Starting over means starting from zero. That friction is the moat.

5. iAM Genie — the Personal AI Layer

Genie is what happens when the LLM stops waiting to be asked. It is the proactive, agentic AI layer that monitors the intent graph, detects patterns in user behaviour, and takes action on the user's behalf — booking, scheduling, emailing, recommending — before the user has to think to ask.

5.1 Reactive mode — answering intent directly

In reactive mode, Genie responds to explicit requests in profile prompts, Stage queries, and direct inputs. The user asks a question. Genie executes the appropriate action and returns a result. The key distinction from a chatbot: Genie does not just answer. It acts.

"Book an appointment at this clinic"	Genie calls the clinic's Calendly API, checks available slots matching the user's calendar, books the appointment, and adds it to Google Calendar
"What does @Gilson love most?"	RAG retrieval over Gilson's intent graph — relevant preference fragments retrieved and fed to LLM for a grounded, personalised answer
"Find me a logo designer under \$200"	Genie queries freelance platform APIs, filters by price and portfolio signals relevant to the user's taste history, surfaces top matches
"Show me only red shoes in my size"	LLM extracts colour and size as filter params, calls the brand's product API or indexed catalogue, returns matching results
"Send the gift link to Sam via WhatsApp"	Genie generates the gift link, formats a message, and triggers the messaging API — no copy-paste required

5.2 Proactive mode — acting before being asked

Proactive Genie monitors the intent graph for behavioural patterns. When a pattern crosses a confidence threshold, Genie surfaces a suggestion — and can complete the action with a single user confirmation.

- User searches 'tooth pain home remedies' three times across two sessions → Genie: "Here are the top-rated dentists near you. Want me to check availability?"
- User has booked Paris trips twice in March → Genie: "It's coming up to your usual Paris window. Want me to start looking at flights?"
- User's favourite restaurant now has Friday evening availability → Genie: "[Restaurant] has a table free this Friday at 7:30. Should I reserve it?"
- User browses running gear but doesn't book → Genie: "The shoes you looked at are on sale today. Want me to add them to your cart?"

5.3 How Genie is built — technically

Step 1 — Pattern or request detection

The intent graph records every interaction with timestamps. A background process (cron or event-driven) runs analysis over the graph looking for repeated query patterns, browse-without-book signals, seasonal recurrence, and proximity to past outcomes. When a signal exceeds the confidence threshold, a Genie event is fired.

Step 2 — Intent extraction and action classification

The Genie event is passed to the LLM with the relevant graph fragments as context. The LLM determines: what is the user trying to accomplish? What type of action does this require? Which API is the correct executor? The output is a structured action plan — not natural language, a typed intent-to-action specification.

Step 3 — RAG retrieval for personalised context

Before the API call fires, the system runs a RAG query over the user's intent graph to retrieve the most relevant contextual fragments: preferred time slots, price sensitivity, brand preferences, past outcomes for similar actions. These fragments are included in the API call parameters where applicable and in any LLM-generated message accompanying the action.

Step 4 — API execution and confirmation

The selected API is called server-side with the constructed parameters. The result is returned to the user in plain language — not a raw API response. If the action requires user confirmation (a booking, a payment), Genie presents the details and waits. One tap confirms. The intent graph is updated with the outcome.

The key principle: Genie has access to the same third-party API layer that powers the Stage. It is not a separate system. It is the Stage with autonomy — the same integrations, called on the user's behalf without them needing to build the prompt.

6. The Profile System

Every iAM entity — user or brand — has a Stage profile. Not a static bio page. A live, queryable, AI-powered presence that responds to any visitor's question using the profile owner's intent graph as its brain.

6.1 User profiles

A user's Stage profile is their public identity and their personal AI presence on the platform. It combines the visual language of Instagram (header image, avatar, bio) with an interactive intelligence layer that no social platform has built: a prompt window that lets any visitor ask anything about the user — and get a genuinely personalised answer.

Header image	Full-width cover — uploaded by the user
Avatar	Profile photo — circular, prominent
Display name	Full name
@handle	Unique identifier — used for tagging in any Stage prompt
Short bio	Up to 160 characters — plain text
Location	City or region — optional
Prompt window	Persistent text input at the bottom — any visitor can query the profile owner's intent graph
Intent graph	Powers all answers — seeded by uploaded document, grown by every platform interaction

What a visitor can ask on a user profile

- "What's Nekha's favourite colour?" → Genie retrieves from Nekha's intent graph via RAG and answers with her documented or inferred colour preferences
- "What kind of restaurants does Nekha like?" → Genie surfaces Nekha's cuisine preferences and past dining history
- "What would Nekha want as a gift?" → Genie cross-references Nekha's wish signals, taste history, and spending comfort to suggest appropriate options
- "Has Nekha been to Tokyo?" → Genie checks her destination history and returns a grounded, factual answer

How a user seeds their profile

At onboarding, iAM provides a downloadable sample document — PDF and DOCX — that guides the user through what to share: interests, travel preferences, dietary choices, favourite items, style signals, colours, brands they love. The user fills it in and uploads it. iAM parses the document, maps each field to the intent graph schema, and the profile is live. The graph continues to grow from every platform action thereafter.

6.2 Brand profiles

A brand's Stage profile is its storefront, its customer service layer, and its AI-powered sales assistant — all in one. Any brand can create a presence on iAM and make it queryable. Visitors do not browse a static product catalogue. They ask, and the brand's intent graph — seeded with product data, brand values, and service offerings — answers.

Header image	Brand cover — full-width
Logo / avatar	Brand logo — circular
Brand name	@handle
Short bio	Brand positioning — 160 characters
Location	HQ city or global — optional
Prompt window	Visitors query the brand's catalogue, services, and offerings in plain English
Brand intent graph	Seeded by brand document — product catalogue, values, target customer, services, pricing
Genie layer	Can take actions on behalf of the visitor — filter products, check availability, book appointments

What a visitor can do on a brand profile

- "Show me only shoes in red, size 7" — inside Nike's Stage profile → Genie extracts colour and size, calls Nike's product API, returns only matching items. No navigation. No filters. One sentence.
- "What's the cheapest flight you have to Dubai next weekend?" — inside an airline's Stage profile → Genie queries the airline's availability API with the extracted parameters and returns the best match

- "Can I get a logo designed for under £300 and how long will it take?" — inside a design agency's profile → Genie answers from the brand's uploaded pricing and delivery document
- "Book a table for two this Friday at 8pm" — inside a restaurant's Stage profile → Genie calls the restaurant's booking API directly and confirms the reservation
- "Do you have a plumber available tomorrow morning?" — inside a home services brand → Genie checks availability via the service's scheduling API and presents open slots

The brand profile principle: Every brand on iAM effectively gets a natural language API on top of their existing systems. Visitors do not need to learn how the brand organises its catalogue. They ask in plain English. Genie does the rest.

Brand document seeding

iAM provides brands with a structured template covering: brand values, product or service catalogue, pricing tiers, typical customer profiles, operational hours, booking or purchasing flow, and FAQs. The uploaded document is parsed and indexed. Brands can re-upload at any time. For larger catalogues, iAM supports direct API connections that keep the brand intent graph continuously updated.

7. Monetization — Built Around the User

Every major internet platform monetises by selling user attention — to advertisers, to the highest bidder, to whoever pays for placement regardless of relevance. iAM is built on the opposite principle. We earn when the user succeeds. Not before. Not instead.

7.1 The problem with how the internet currently makes money

The internet's dominant business model is fundamentally at odds with the user's interest. Advertising platforms earn on impressions and clicks — whether or not a purchase ever happens. This creates a structural incentive to maximise attention at the expense of outcomes. Feeds are engineered for doom-scrolling. Search results are padded with ads before organic content. Recommendation algorithms optimise for engagement, not satisfaction.

Google	Charges per click — whether or not a purchase happens. Revenue is decoupled from the user's outcome.
---------------	--

Meta	Charges per impression — whether or not anyone cares. Revenue scales with time-in-app, not with user value received.
-------------	--

Amazon	Charges for sponsored placement — products that paid rank above products that fit. Relevance is purchasable.
iAM	Charges nothing until the intent becomes a completed purchase. Revenue and user outcome are the same event.

7.2 iAM's revenue model

Core: transaction percentage

In the early period and throughout, iAM's primary revenue is a percentage of every booking or purchase completed through the Stage. The rate varies by category and vendor relationship. This model is inherently aligned: iAM only makes money when the user completes an action they wanted to complete. There is no revenue from browsing. No revenue from clicking. No revenue from simply seeing something.

Post-traction: sponsored placements — on our terms

Once the Intent Graph has sufficient signal to enforce relevance with integrity, iAM will introduce sponsored placements. These are categorically different from advertising as it exists today.

- A sponsored card is a result that has already passed every relevance and quality filter for that intent and that user. The brand has chosen to bid for a higher position within the already-approved, already-relevant set. If it does not genuinely fit the intent, it does not enter at any price.
- A sponsored row is a brand committing specific products to own an entire row within qualifying Stages — but every product in that row passes the same relevance gate as everything else. Nike can sponsor the 'What to wear' row for Paris trip Stages. Not all Stages. Only those where their products genuinely belong.
- Sponsored elements are always clearly labelled. The word 'Sponsored' is part of the interface — not fine print, not a tooltip. One word, always visible.
- Brands pay on completed purchase only — not for impressions, not for placement. A brand bets on the outcome. iAM only earns when they earn.

7.3 What iAM will never do

These lines do not move. They are not policy. They are architecture.

- iAM will never show a result only because someone paid. Payment shifts position within the qualified set. It cannot create qualification.
- iAM will never sell user data. The Intent Graph belongs to the user. It is used exclusively to improve their Stage experience. It is not a data asset for advertisers.

- iAM will never engineer for engagement at the expense of outcome. The feed does not exist. The Stage exists only when intent exists. There is no mindless scroll surface.
- iAM will never run banners, pop-ups, interstitials, or pre-rolls. The only valid commercial presence inside a Stage is a card or row that honestly answers the user's intent.
- iAM will never blur the line between organic and sponsored results. The distinction is always visible, always clear, and always honest.

iAM does not sell your attention. It earns on your outcome.

7.4 Why this model works financially

iAM's transaction-only model is not a compromise. It is a competitive advantage. Users trust the Stage because they know nothing in it exists to manipulate them. That trust produces higher engagement, higher booking rates, and higher average order values than any attention-capture platform can achieve. A user who trusts the result converts at a fundamentally higher rate than a user who suspects they are being sold to.

For brands, sponsored placement in iAM is lower risk than any alternative. Google charges for the click whether the sale happens or not. Meta charges for the impression regardless. In iAM, the brand pays only when the user buys. For brands, that is the most efficient marketing spend available.

8. Where We Are and Where We Are Going

8.1 Right now — the travel vertical

iAM launches with travel as its proving ground. Travel is the highest-intent, highest-value, most coordination-heavy category of human need. Planning a trip currently requires the most tabs, the most platforms, the most back-and-forth. Solving it completely — flights, hotels, cars, experiences, restaurants, weather, maps, gifting, multi-person checkout — proves the core architecture at its hardest. If iAM can make travel effortless, it can make everything else effortless.

FEATURE	PROVIDER	STATUS AT LAUNCH
---------	----------	------------------

Flights	Duffel	Real-time search, offer selection, booking confirmation
Hotels & stays	Duffel	Filtered by group size, dates, and budget signal
Rental cars	Duffel	Category, passenger count, booking
Activities & tours	Viator	Search, availability, booking
Restaurant bookings	OpenTable	Availability, party size, reservation creation
Weather forecast	OpenWeatherMap	Forecast for destination dates
Maps & discovery	Google Maps	Nearby POIs, shopping, landmarks
Gift Link	Stripe + Vendor	Address-decoupled gifting across all bookable items
Collaborative Stage	SSE + Redis	Real-time multi-participant planning and checkout
User & brand profiles	MongoDB + LLM	Queryable public profiles with intent graph answers

8.2 The expansion roadmap — everything a person needs

Travel is the wedge. The architecture is built for everything. Every new vertical iAM enters follows the same pattern: identify the APIs that cover that domain, connect them to the Stage, extend the intent parsing logic to understand domain-specific language, and let the Intent Graph learn user preferences in that category. The platform does not need to be rebuilt. It needs to be extended.

VERTICAL	HORIZON	WHAT IT ADDS
Shopping	Near-term	Direct product search and purchase across all major retailers, fashion, electronics, groceries — searchable by natural language across every store at once
Home services	Near-term	Plumbers, electricians, cleaners, handymen, gardeners — Genie checks availability, books, confirms. 'I need a plumber tomorrow morning' is one sentence.
Digital services	Near-term	Logo designers, web developers, copywriters, photographers, domain registration, hosting — the full stack of what a person or business needs to build something online

Health & wellness	Near-term	Doctors, dentists, physiotherapists, therapists, gyms — Genie books appointments, tracks follow-ups, and surfaces proactive health suggestions from intent patterns
Education	Mid-term	Courses, tutors, certifications, workshops — assembled around a learning intent. 'I want to learn Python in 8 weeks' produces a full curriculum Stage with bookable sessions.
Local & food	Mid-term	Grocery delivery, local restaurant ordering, neighbourhood services — everything hyperlocal, assembled in the Stage and acted on immediately
iAM Wallet	Mid-term	An in-app wallet that removes Stripe's checkout flow entirely. Users top up once. Every purchase inside iAM deducts from the wallet instantly. Split bills, group contributions, and gift budgets managed natively.
iAM Messaging	Mid-term	Direct messaging between users — sharing Stages, collaborating on plans, sending gift links, building trips together. Communication native to intent, not a separate app.
iAM Social	Long-term	One platform to hold all connected social contexts. Follow people, share Stages publicly, see what your network is planning, discover new places and products through trusted social graphs — not algorithmic feeds.
Finance & planning	Long-term	Smart budgeting powered by the Intent Graph. Spend analytics, savings goals, travel budgets — the graph already tracks what you spend and where. Finance is a natural extension.
iAM Health	Long-term	A full personal health management layer. Prescription reminders, medical appointment tracking, health record access, proactive Genie suggestions based on health intent patterns.
B2B services	Long-term	Everything a business needs: legal, accounting, HR, logistics, marketing. The Stage serves business intent the same way it serves personal intent — one prompt, every relevant service assembled.

The expansion principle: iAM does not need a different architecture for each new vertical. It needs a new API integration, extended intent parsing, and a new category of cards. Every vertical iAM enters makes every other vertical richer — because the Intent Graph learns across all of them.

9. iAM as a Super-App

A super-app is not a platform with many features. It is a platform that becomes indispensable — the one app a person reaches for regardless of what they need. WeChat did it in China. Grab did it in Southeast Asia. Neither exists as a global, AI-native, intent-first platform. That is what iAM is building.

9.1 Why the Intent Graph makes iAM different from every other super-app attempt

Every super-app attempt in Western markets has failed because they tried to aggregate features. They added a wallet to a chat app. They added shopping to a social feed. The features did not talk to each other. The platform did not learn. Users did not feel like the platform knew them — because it did not.

iAM is different because the Intent Graph cuts across every vertical. A user's food preferences inform their restaurant suggestions, their grocery delivery, their travel dining choices, and their Genie suggestions — simultaneously. Their spending signal shapes hotel recommendations, product results, and digital service options. Their calendar awareness means Genie suggests the dentist during a light week and the restaurant on a Friday evening. Every feature on the platform is smarter because of every other feature.

9.2 Network effects that compound

- Every user who joins enriches their own intent graph — their Stage improves with every action
- Every brand that joins makes user Stages richer — more products, more services, more answers to profile queries
- Every booking outcome improves ranking models for all similar users — the platform learns collectively
- Social features (messaging, shared Stages, public profiles) add a network dimension — each new user makes existing users' experience more connected
- iAM Wallet makes every transaction faster — which increases booking rates — which generates more outcome signal — which improves ranking — which increases booking rates

Switching cost at scale: A user with two years of Intent Graph history on iAM has a personalised AI layer that knows more about their preferences than any other system ever has. Moving to a competitor means explaining yourself again from scratch. That is a switching cost no price or feature can easily overcome.

10. We Are the New Internet

The internet was built to give people access to information. iAM is built to take action on it.

For thirty years, the internet has been a place you visit to find things. You search. You browse. You decide. You navigate. You pay. You do it again on another platform. You carry all the cognitive load of connecting the dots between what you want and what the internet can offer.

iAM ends that loop. One prompt. The internet answers. Every relevant service, product, provider, and piece of content assembled around the user's specific intent and their specific history. Collaborative with anyone they choose. Actionable in one checkout. Remembered for next time. Proactively suggested before they even have to ask.

iAM is not a better search engine. It is not a smarter recommendation algorithm. It is not another platform competing for your attention. It is the interface layer that finally sits between the user and the internet and works for the user — not for advertisers, not for engagement metrics, not for the highest bidder.

The old internet

You search. You browse. You compare. You tab-switch. You copy-paste. You pay on four platforms. You repeat.

iAM

You say what you want. iAM assembles the internet around you. You act once. Done.

We Are the New Internet.

iAM · Intent becomes action.